



```

root@localhost:/home/training
File Edit View Search Terminal Help
OK
Time taken: 0.01 seconds
hive> create table twitter_data (label int ,number int,datee string, no_query st
OK
Time taken: 0.284 seconds
hive> load data local inpath '/home/training/t_data.csv' into table twitter_data
Copying data from file:/home/training/t_data.csv
Copying file: file:/home/training/t_data.csv
Loading data to table twitter.twitter_data
OK
Time taken: 0.261 seconds
hive> select * from twitter_data limit 5;
OK
0      1467810369      Mon Apr 06 22:19:45 PDT 2009      NO_QUERY      _TheSpec
0      1467810672      Mon Apr 06 22:19:49 PDT 2009      NO_QUERY      scottham
0      1467810917      Mon Apr 06 22:19:53 PDT 2009      NO_QUERY      mattycus
0      1467811184      Mon Apr 06 22:19:57 PDT 2009      NO_QUERY      ElleCTF
0      1467811193      Mon Apr 06 22:19:57 PDT 2009      NO_QUERY      Karoli
Time taken: 0.196 seconds
hive> create table dyn_part(number int,datee string,no_query string,name string,
OK
Time taken: 0.033 seconds
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;

```

```

root@localhost:/home/training
File Edit View Search Terminal Help
Time taken: 0.033 seconds
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> insert into table dyn_part partition(label) select number,datee,no_query,n
Total MapReduce jobs = 2
Launching Job 1 out of 2
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1704618892832_0006, Tracking URL = http://localhost.localdoma
Kill Command = /usr/lib/hadoop/bin/hadoop job -Dmapred.job.tracker=localhost:91
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-01-07 04:52:11,101 Stage-1 map = 0%, reduce = 0%
2024-01-07 04:52:14,342 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.63 se
MapReduce Total cumulative CPU time: 630 msec
Ended Job = job_1704618892832_0006
Ended Job = 1761817844, job is filtered out (removed at runtime).
Moving data to: hdfs://localhost:8020/tmp/hive-root/hive_2024-01-07_04-52-05_800
Loading data to table twitter.dyn_part partition (label=null)
Loading partition {label=__HIVE_DEFAULT_PARTITION__}
Loading partition {label=4}
Loading partition {label=0}
22011 Rows loaded to dyn_part
MapReduce Jobs Launched:
Job 0: Map: 1 Cumulative CPU: 0.63 sec HDFS Read: 3037077 HDFS Write: 270104
Total MapReduce CPU Time Spent: 630 msec

```

Applications

File Edit View VM Tabs Help

Cloudera-Udacity Training...

Sun Jan 7, 4:53 AM

File Edit View History go Send email

localhost:50070/dfshealth.jsp

The Platform for Big Data

NameNode 'localhost.localdomain:8020' (active)

Started: Sun Jan 07 04:14:27 EST 2024

Version: 2.0.0-cdh4.1.1, 581959ba23e4af85afdbdb98b7687662fe9c5f20

Compiled: Tue Oct 16 11:07:59 PDT 2012 by jenkins from Unknown

Upgrades: There are no upgrades in progress.

Cluster ID: CID-5920c32c-c3f1-42e5-81a0-c93032e8442a

Block Pool ID: BP-2100437999-127.0.0.1-1355878275444

Browse the filesystem

NameNode Logs

Cluster Summary

Security is OFF

581 files and directories, 278 blocks = 859 total.

Heap Memory used 11.1 MB is 51% of Committed Heap Memory 21.58 MB. Max Heap Memory is 193.38 MB.

Non Heap Memory used 24.14 MB is 99% of Committed Non Heap Memory 24.38 MB. Max Non Heap Memory is 96 MB.

Configured Capacity: 14.08 GB

Dfs Used: 51.24 MB

Non Dfs Used: 4.31 GB

Dfs Remaining: 9.72 GB

Dfs Used%: 0.36 %

Dfs Remaining%: 69.01 %

Block Pool Used: 51.24 MB

Block Pool Used%: 0.36 %

DataNodes usages: Min % 0.36 % Median % 0.36 % Max % 0.36 % stdev % 0 %

Live Nodes: 1 (Decommissioned: 0)

Dead Nodes: 0 (Decommissioned: 0)

Decommissioning Nodes: 0

Number of Under-Replicated Blocks: 11

NameNode Journal Status:

Current transaction ID: 12594

Journal Manager

State

Applications

File Edit View VM Tabs Help

Cloudera-Udacity Training...

Sun Jan 7, 4:53 AM

File Edit View History bookmarks Tools Help

localhost:50070/browseDirectory.jsp?nameNodeUrl=50070&nameNode=localhost.localdomain:8020

The Platform for Big Data

Contents of directory /

Go to go

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
hbase	dir				2023-12-20 02:20	rw-rw-r--	hbase	supergroup
sample	dir				2024-01-02 14:21	rw-rw-r--	root	supergroup
tmp	dir				2023-12-20 02:44	rw-rw-rw-	hdfs	supergroup
test	dir				2024-01-02 14:31	rw-rw-r--	root	supergroup
user	dir				2023-12-19 08:57	rw-rw-rw-	hadoop	supergroup
var	dir				2014-07-27 09:23	rw-rw-r--	hdfs	supergroup

Go back to DFS home

Local logs

Log directory

Hadoop 2024

localhost:50070/dfshealth.jsp

Applications Places System

HDFS:user - Mozilla Firefox

Sun Jan 7, 4:53 AM

File Edit View History Bookmarks Tools Help

HDFS:User

localhost.localdomain:50075/browseDirectory.jsp?ds=%2Fuser%2Fhive&nameNodeInfoPort=50076&nameNode=localhost.localdomain:8020

The Platform for Big Data

Hadoop JobTracker

Hadoop NameNode

Contents of directory /user

Go to: /user

go

Go to parent directory

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
history	dir				2014-07-27 09:35	rxwxrwxrwx	yarn	supergroup
hive	dir				2013-09-05 20:08	rxwxr-xr-x	hue	supergroup
hue	dir				2013-09-30 18:37	rxwxr-xr-x	hue	hue
root	dir				2013-12-19 00:57	rxwx---	root	supergroup
training	dir				2014-07-27 09:44	rxwxr-xr-x	training	supergroup
trial	dir				2014-07-27 03:52	rxwxr-xr-x	training	supergroup

Go back to DFS home

Local logs

Log directory

Hadoop, 2024.

root@localhost:/home/...

HDFS:user - Mozilla Firefox

root@localhost:/home/...

File Edit View History Bookmarks Tools Help

HDFS:User/hive

localhost.localdomain:50075/browseDirectory.jsp?ds=%2Fuser%2Fhive&nameNodeInfoPort=50076&nameNode=localhost.localdomain:8020

The Platform for Big Data

Hadoop JobTracker

Hadoop NameNode

Contents of directory /user/hive

Go to: /user/hive

go

Go to parent directory

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
warehouse	dir				2024-01-07 04:48	rxwxrwxrwx	hue	supergroup

Go back to DFS home

Local logs

Log directory

Hadoop, 2024.

root@localhost:/home/...

HDFS:user/hive - Mozilla Firefox

root@localhost:/home/...



Applications

File Edit View VM Tabs Help

Home Cloudera-Udacity Training...

Sun Jan 7, 4:53 AM

File Edit View History Bookmarks Tools Help

HDFS/user/hive/warehouse

localhost.localdomain:50075/browse/Directory.jsp?ds=%2Fuser%2Fhive%2Fwarehouse&nameNodeHofPort=50070&readw=localhost.localdomain:8020

The Platform for B... Hadoop JobTracker Hadoop NameNode

Contents of directory [user/hive/warehouse](#)

Go to:

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permissions	Owner	Group
demo.db	dir				2023-12-21 00:37	rwxr-xr-x	root	supergroup
lrldb	dir				2023-12-20 02:31	rwxr-xr-x	training	supergroup
iris_data.db	dir				2023-12-20 02:40	rwxr-xr-x	training	supergroup
supply_chain.db	dir				2024-01-06 14:57	rwxr-xr-x	root	supergroup
twitter.db	dir				2024-01-07 04:50	rwxr-xr-x	root	supergroup
twitter_data	dir				2024-01-07 04:35	rwxr-xr-x	root	supergroup

[Go back to DFS home](#)

Local logs

[Log directory](#)

Hadoop, 2024.

Applications

File Edit View VM Tabs Help

Home Cloudera-Udacity Training...

Sun Jan 7, 4:53 AM

File Edit View History Bookmarks Tools Help

HDFS/user/hive/warehouse

localhost.localdomain:50075/browse/Directory.jsp?ds=%2Fuser%2Fhive%2Fwarehouse%2Ftwitter.db&nameNodeHofPort=50070&readw=localhost.localdomain:8020

The Platform for B... Hadoop JobTracker Hadoop NameNode

Contents of directory [user/hive/warehouse/twitter.db](#)

Go to:

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permissions	Owner	Group
dyn_part	dir				2024-01-07 04:52	rwxr-xr-x	root	supergroup
twitter_data	dir				2024-01-07 04:49	rwxr-xr-x	root	supergroup

[Go back to DFS home](#)

Local logs

[Log directory](#)

Hadoop, 2024.

Applications | File | Edit | View | VM | Tabs | Help

Cloudera-Udacity Training...

Sun Jan 7, 4:54 AM

file | edit | view | history | bookmarks | tools | help

HDFS/userhive/warehouse/...

localhost.localdomain:5007/browseDirectory.jsp?dir=%2Fuser%2Fhive%2Fwarehouse%2Ftwitter.db%2Fdyn\_part%2Fusernode%2Fpart=50070&read=localhost.localdomain:8020

The Platform for Big Data | Hadoop JobTracker | Hadoop NameNode

Contents of directory [/user/hive/warehouse/twitter.db/dyn\\_part](#)

Go to [parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
<a href="#">label-0</a>	dir				2024-01-07 04:52	rwxr-xr-x	root	supergroup
<a href="#">label-4</a>	dir				2024-01-07 04:52	rwxr-xr-x	root	supergroup
<a href="#">label-HIVE_DEFAULT_PARTITION_</a>	dir				2024-01-07 04:52	rwxr-xr-x	root	supergroup

[Go back to DFS home](#)

Local logs

[Log directory](#)

[Hadoop](#), 2024.

root@localhost: /home... | HDFS/userhive/ware... | root@localhost: /home...

Applications | File | Edit | View | VM | Tabs | Help

Cloudera-Udacity Training...

Sun Jan 7, 4:54 AM

file | edit | view | history | bookmarks | tools | help

HDFS/userhive/warehouse/...

localhost.localdomain:5007/browseDirectory.jsp?dir=%2Fuser%2Fhive%2Fwarehouse%2Ftwitter.db%2Fdyn\_part%2Flabel%3D0&usernode%2Fpart=50070&read=localhost.localdomain:8020

The Platform for Big Data | Hadoop JobTracker | Hadoop NameNode

Contents of directory [/user/hive/warehouse/twitter.db/dyn\\_part/label-0](#)

Go to [parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
<a href="#">000000.0</a>	file	1.47 MB	1	64 MB	2024-01-07 04:52	rw-r--r--	root	supergroup

[Go back to DFS home](#)

Local logs

[Log directory](#)

[Hadoop](#), 2024.

root@localhost: /home... | HDFS/userhive/ware... | root@localhost: /home...

Applications | File | Edit | View | VM | Tabs | Help

Cloudera-Udacity Training...

Sun Jan 7, 4:54 AM

Applications

File Edit View VM Tabs Help

Cloudera Udacity Training...

Sun Jan 7, 4:54 AM

File Edit View History Bookmarks Tools Help

localhost.localdomain:50075/browse/Directory.jsp?dir=%2Fuser%2Fhive%2Fwarehouse%2Ftwitter.db%2Fdyn\_part%2Flabel%3D0&name=node%2Fport=50075&rsadd=localhost.localdomain:8020

The Platform for Bi... Hadoop JobTracker Hadoop NameNode

Contents of directory /user/hive/warehouse/twitter.db/dyn\_part/label=0

Goto: /user/hive/warehouse/twitter.db/dyn\_part/label=0 go

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
000000_0	file	1.47 MB	1	64 MB	2024-01-07 04:52	rw-r--r--	root	supergroup

[Go back to DFS home](#)

### Local logs

[Log directory](#)

Hadoop, 2024.

root@localhost: /home... HDFS: /user/hive/warehouse/... root@localhost: /home...

The Platform for Bi... Hadoop JobTracker Hadoop NameNode

Contents of directory /user/hive/warehouse/twitter.db/dyn\_part/label=0

Goto: /user/hive/warehouse/twitter.db/dyn\_part/label=0 go

[Go to parent directory](#)

Name	Type	Size	Replication	Block Size	Modification Time	Permission	Owner	Group
000000_0	file	1.47 MB	1	64 MB	2024-01-07 04:52	rw-r--r--	root	supergroup

[Go back to DFS home](#)

### Local logs

[Log directory](#)

Hadoop, 2024.